RESEARCH PROPOSAL

# Architecting a Secure, Compliant, and Intelligent CBDC Ecosystem: A Cryptographic and Agentic Governance Framework for AI–Blockchain Convergence

*Bridging the Privacy–Compliance Paradox, the Algorithmic Liability Void, and the Determinism–Adaptability Trade-off in Central Bank Digital Currency Infrastructure*

---

**Domain:** Applied Machine Learning, Financial Cryptography, Distributed Ledger Technology, Regulatory Technology (RegTech)

**Proposal Type:** Multi-phase academic research proposal with parallel patent-positioning track

**Target Readiness:** Progression toward Technology Readiness Level 6 (TRL-6) — system/subsystem model or prototype demonstration in a relevant environment

## 1. Abstract

Over one hundred central banks, representing more than 98 percent of global GDP, are actively exploring or deploying Central Bank Digital Currencies (CBDCs), while global AI spending is projected to reach USD 2.5 trillion by 2026 and tokenized assets are projected to reach double-digit trillions by 2030. The convergence of AI, blockchain, and CBDCs exposes three unresolved structural contradictions: a privacy-compliance paradox in AML/CFT enforcement, a legal and liability void for autonomous algorithmic decisions, and an architectural trade-off between blockchain's deterministic consensus and AI's probabilistic adaptability. This proposal sets out a three-phase research program to decouple AI execution from blockchain consensus while preserving mathematical verifiability, legal accountability, and quantum-era resilience. The program integrates Zero-Knowledge Machine Learning (ZKML) and Optimistic ZK-ML (zk-OPML) for verifiable off-chain inference, Federated Learning (FL) for privacy-preserving cross-border AML modelling, AgentBound Tokens (ABTs) for cryptoeconomic liability on autonomous financial agents, and lattice-based post-quantum zero-knowledge proofs for long-horizon cryptographic security. Each phase is structured around falsifiable hypotheses with explicit failure criteria, and the resulting system architecture, cryptographic protocols, and governance token design are positioned for both peer-reviewed publication and patent filing.

## 2. Introduction and Background

The modernization of the global financial system is being propelled by the convergence of three previously distinct technological paradigms: Central Bank Digital Currencies, distributed ledger technology, and Artificial Intelligence. Blockchain has matured beyond cryptocurrency into a foundational infrastructure for coordinating trust, verifiable ownership, and programmable incentives, while AI — through generative models, machine learning, and autonomous agentic systems — has become indispensable for real-time fraud detection, dynamic risk management, and automated regulatory compliance.

This convergence is already restructuring physical and economic capital: compute facilities once dedicated to cryptocurrency mining are being repurposed for AI cloud infrastructure, and a new class of multi-function crypto assets is emerging that sits between utility, security, and decentralized compute provision, raising unresolved classification questions for regulators. Despite this economic momentum, the integration of AI, blockchain, and CBDCs is fraught with structural, regulatory, and technological gaps that current systems are not designed to resolve.

## 3. Problem Statement

Current digital currency infrastructures cannot securely balance the public's right to financial privacy with the state's mandate for Anti-Money Laundering (AML) and Counter-Terrorism Financing (CFT) compliance.

Blockchain's transparent ledger structure allows AI-driven chain-analysis tools to deanonymize users through transaction graph clustering, while privacy-enhancing technologies applied to counter this severely hamper the detection of illicit financial flows. Compounding this, AI-based AML systems face severe label scarcity: in the widely used Elliptic Bitcoin dataset, 77 percent of transactions are unlabeled and only 2 percent are confirmed illicit, biasing supervised models toward the majority licit class.

As autonomous AI agents increasingly execute financial operations on behalf of delegated principals, existing legal frameworks — which rest on human intent, negligence, and direct contractual agreement — fail to assign liability for algorithmic errors or opaque "black box" decisions. Empirical analysis of early AI-CBDC pilots indicates a persistent 15–20 percent deficit in consumers successfully obtaining remedies for AI-driven false positives, disproportionately affecting small and rural merchants. There is therefore a critical need to decouple AI execution from blockchain consensus while ensuring mathematical verifiability, legal accountability, and resilience against future cryptographic threats, including quantum computing attacks on the elliptic-curve cryptography and hash functions that currently secure CBDC infrastructure.

## 4. Research Gaps Addressed

### 4.1 The Privacy–Compliance Paradox

A functional gap exists in which privacy-preserving mechanisms blind AI-based compliance systems, while AI-based compliance mechanisms erode the privacy guarantees CBDCs are expected to provide. No production-grade architecture currently resolves this without a centralized trusted intermediary.

### 4.2 The Algorithmic Accountability Deficit

No existing legal or technical framework assigns enforceable, automatically executable liability to an autonomous AI agent operating within a CBDC network. Accountability is diffused across developers, data suppliers, wallet-operating intermediaries, and issuing central banks, with no cryptoeconomic mechanism to internalize the risk of agent error.

### 4.3 The Determinism–Adaptability Trade-off

Smart contracts require absolute determinism for consensus, while the AI models needed to interpret unstructured data, assess dynamic credit risk, or model macroeconomic conditions are inherently non-deterministic and computationally infeasible to execute on-chain. A verifiable bridge between the two paradigms — without reintroducing a centralized oracle as a single point of failure — remains an open research problem.

## 5. Research Objectives

1. Intelligence-Driven Monetary Policy: Develop algorithmic monetary policy engines using Deep Reinforcement Learning (DRL) for dynamic, targeted macroeconomic interventions based on high-frequency, privacy-preserved CBDC transaction data.
2. Cryptographic Compliance and Privacy: Design and empirically validate Zero-Knowledge Machine Learning (ZKML) and Federated Learning (FL) frameworks that allow AI models to assess AML/CFT risk without exposing raw personally identifiable information (PII).
3. Agentic Governance and Quantum Resilience: Design a governance framework that enforces cryptoeconomic liability on autonomous AI agents through AgentBound Tokens (ABTs), while upgrading the underlying cryptographic architecture to withstand quantum computing attacks.

## 6. Falsifiable Hypotheses and Failure Criteria

Consistent with a novelty-first, evidence-driven research design, each objective is decomposed into a falsifiable hypothesis with an explicit, pre-registered failure criterion, so that a negative result is scientifically informative rather than merely inconclusive.

### *H1 — Verifiable Off-Chain Inference (zk-OPML)*

**Hypothesis:** Restructuring AML/credit-risk inference as a bisectable ONNX computation graph under an Optimistic-ML dispute protocol (zk-OPML) will reduce end-to-end proof-generation latency and on-chain verification cost by at least an order of magnitude relative to monolithic ZKML proofs, without degrading model accuracy by more than 1 percentage point.

**Failure Criterion:** The hypothesis is rejected if the bisection protocol fails to isolate a disputed operator within a bounded number of interactive rounds on models exceeding a defined depth, or if latency/cost reduction falls below one order of magnitude on the benchmark model suite.

### *H2 — Federated AML Detection Under Extreme Label Scarcity*

**Hypothesis:** A blockchain-orchestrated Federated Learning framework trained across simulated multi-institution data silos will achieve a statistically significant reduction in false-positive AML flags (target: ≥20 percent) relative to a centralized baseline trained on the same label-scarce distribution (e.g., the Elliptic dataset's 2 percent illicit label rate), while formally bounding PII leakage via differential privacy guarantees (target $\varepsilon$).

**Failure Criterion:** The hypothesis is rejected if the federated model's false-positive reduction is not statistically significant at the pre-registered threshold, or if the differential privacy budget required to reach that reduction exceeds an acceptable $\varepsilon$, indicating an unresolved privacy-utility trade-off.

### *H3 — Cryptoeconomic Liability via AgentBound Tokens*

**Hypothesis:** Binding autonomous financial agents to staked, non-transferable AgentBound Tokens (ABTs) with coded slashing conditions will produce a measurable reduction in policy-violating agent actions in a simulated multi-agent CBDC testbed, relative to an unstaked control group of equivalent agents.

**Failure Criterion:** The hypothesis is rejected if staking and slashing produce no statistically significant behavioral deterrence relative to control, or if the slashing mechanism can be circumvented through Sybil-style agent re-identification within the testbed.

***H4 — Post-Quantum Zero-Knowledge Proof Feasibility***

**Hypothesis:** A lattice-based ZK-STARK construction can secure ZKML proof verification for the model classes used in H1–H3 while keeping proof size and verifier latency within a bounded overhead factor relative to classical (non-quantum-resistant) ZK-STARKs, making PQC migration practical without re-architecting the settlement layer.

**Failure Criterion:** The hypothesis is rejected if the overhead factor exceeds the bound required for real-time or near-real-time transaction verification on the target hardware profile.

## 7. Proposed Methodology

### Phase 1 — Bridging Determinism and Privacy (Months 1–8)

This phase implements advanced cryptographic frameworks such as zk-OPML to bridge probabilistic AI and deterministic blockchain consensus. By defaulting to the assumption that off-chain AI inference results are correct, the protocol provides a challenge window during which network participants may dispute a result. If challenged, an interactive bisection protocol isolates the exact neural-network operator causing the disagreement, generating a targeted zero-knowledge proof for that single operator rather than the entire model, sharply reducing transaction cost.

In parallel, a Federated Learning layer will be deployed on a decentralized blockchain orchestration layer, drawing on architectures such as BAML (Blockchain Anti-Money Laundering), so that central and commercial banks can collaboratively train shared AML models on local data without transmitting raw citizen-level PII across institutional or national boundaries. Model updates will be shared as encrypted, differentially private gradients and aggregated via secure multi-party computation.

### Phase 2 — Resolving the Liability Void via Agentic Finance (Months 7–16)

This phase develops and tests AgentBound Tokens (ABTs) — non-transferable cryptographic credentials that bind an autonomous AI agent to an immutable on-chain record of its behavior, training provenance, and compliance history. Developers will be required to stake financial collateral tied to the ABT via a Proof-of-Stake mechanism, with explicit, protocol-enforced slashing conditions that automatically confiscate staked collateral if an agent violates sanctions rules or exhibits discriminatory behavior.

The system will additionally embed verifiable Explainable AI (an ExpProof-style construction) so that AI-generated explanations are cryptographically tied to a committed model, guaranteeing that automated decisions are genuinely auditable and consistent with right-to-explanation obligations rather than post-hoc justifications generated to satisfy an audit.

### Phase 3 — Quantum Security and Real-Time Monetary Policy (Months 15–24)

The final phase formulates hybrid blockchain architectures integrated with Post-Quantum Cryptography (PQC) to safeguard against Shor's and Grover's algorithms, including quantum-resistant, lattice-based ZK-STARK protocols capable of securing large machine-learning parameter sets without sacrificing scalability. The phase concludes by testing the integration of a Reverse Kelly Automated Market Maker (AMM) alongside Deep Reinforcement Learning models that ingest real-time, privacy-preserved transaction data to autonomously optimize liquidity distribution and adjust micro-targeted interest rates, removing the temporal lag inherent in centralized policy committees.

## 8. System Architecture Overview

| Layer | Function | Primary Mechanism |
|---|---|---|
| Settlement (Layer 1) | Sovereign, risk-free atomic settlement for systemically important transactions (Wholesale CBDC). | Central bank–issued wCBDC ledger |
| Credit Creation (Layer 2) | Tokenized deposits within the commercial banking system, preserving credit creation without disintermediation. | Bank-issued tokenized deposits |
| Distribution (Layer 3) | Borderless retail distribution and innovation via privately issued, reserve-backed tokens. | Regulated stablecoins |
| Compliance & Inference (cross-cutting) | Off-chain AML/credit inference with on-chain verifiability and PII protection. | ZKML / zk-OPML, Federated Learning |
| Agentic Governance (cross-cutting) | Binds autonomous agents to staked, auditable, and slashable on-chain identity. | AgentBound Tokens (ABTs) |

## 9. Novelty and Patent Positioning

The research program is designed for dual-track output: peer-reviewed academic publication and patent filing, consistent with a novelty-first framing at each design decision.

### 9.1 Novelty Claims

- A dispute-resolution architecture that combines optimistic ML execution (zk-OPML) with selective, operator-level zero-knowledge proof generation specifically tuned to CBDC-scale AML and credit-inference models, rather than generic ML workloads.
- A cryptoeconomic liability instrument (the AgentBound Token) that binds staking, slashing, and verifiable explanation generation into a single non-transferable on-chain credential scoped to regulated financial agents.

- A federated, blockchain-orchestrated AML training pipeline that jointly optimizes for detection performance under extreme label scarcity and a formally bounded differential-privacy guarantee, evaluated specifically against sovereign CBDC compliance requirements.
- A lattice-based post-quantum verification layer sized specifically for the operator-level proofs produced by the zk-OPML dispute protocol, rather than a general-purpose PQC migration.

### 9.2 Patentability Considerations under the Indian Patents Act, 1970

Because the core contributions combine a specific cryptographic protocol (the bisection-based dispute mechanism and the staking/slashing logic bound to the ABT) with a defined technical effect (reduced proof-generation cost, enforceable agent liability, and bounded privacy leakage), the resulting claims will be drafted to emphasize a concrete technical implementation and a measurable improvement in system performance or security, rather than an abstract algorithm or business method standing alone. This distinction is central to satisfying the exclusion under Section 3(k) of the Patents Act, 1970, which bars patents for algorithms and business methods "per se" but does not bar inventions that apply an algorithm to produce a specific, verifiable technical effect within a described system architecture. Patent drafting will proceed in parallel with the academic manuscript, with claims scoped around: (i) the operator-level bisection and selective proof-generation method, (ii) the ABT staking/slashing state-transition logic, and (iii) the specific federated-aggregation protocol used for cross-border AML training. A formal patentability and prior-art search is recommended as an early milestone in Phase 1 to finalize claim scope before public disclosure of the corresponding manuscript.

## 10. Expected Outcomes and Significance

The research is expected to yield a three-layer monetary architecture that supports a fast-paced, autonomous agent economy while satisfying sovereign compliance requirements. By combining compliance-by-design mechanisms with mathematical verifiability, the resulting CBDC ecosystem architecture is intended to support straight-through processing for cross-border transactions, drawing on the compliance-embedding approach demonstrated by BIS initiatives such as Project Mandala and Project Dunbar. The program's central contribution is to transform three currently unresolved technological gaps — the privacy-compliance paradox, the algorithmic liability void, and the determinism-adaptability trade-off — into systemic architectural strengths, aligning AI's analytical adaptability with blockchain's cryptographic trust guarantees.

Beyond the immediate technical contributions, the program is expected to inform policy discussion on "Know Your Agent" (KYA) frameworks as a successor to Know Your Customer (KYC) in an agentic financial system, and to produce a reusable evaluation methodology — grounded in the pre-registered hypotheses in Section 6 — for assessing verifiable-AI proposals in future CBDC pilot programs.

## 11. Technology Readiness and Milestone Roadmap

| Months | Phase | Key Milestone | Target TRL |
|---|---|---|---|
| 1–8 | Phase 1: Determinism & Privacy | zk-OPML bisection prototype; federated AML pipeline benchmark (H1, H2) | TRL 3–4 |
| 7–16 | Phase 2: Agentic Liability | ABT staking/slashing testbed; ExpProof-style explanation audit (H3) | TRL 4–5 |
| 15–24 | Phase 3: Quantum Security & Policy | Lattice-based ZK-STARK integration; DRL policy-engine simulation (H4) | TRL 5–6 |
| 22–24 | Consolidation | System-level integration demonstration in a relevant simulated CBDC environment; manuscript and patent filing | TRL 6 |

## 12. Selected References

*A full bibliography of primary sources — including BIS project reports (Mandala, Dunbar, Polaris), IMF working papers, and the ZKML/zk-OPML/zkComposer/ExpProof/ABT literature underlying this proposal — is maintained separately and available on request. Representative sources include:*


- Bank for International Settlements, Project Mandala: Streamlining cross-border transaction compliance.
- Bank for International Settlements, Project Dunbar / hybrid CBDC settlement models.
- ExpProof: Operationalizing Explanations for Confidential Models with ZKPs, arXiv:2502.03773.
- zk-OPML: Using zero-knowledge proofs to optimize OPML, ResearchGate.
- zkComposer: Decomposing Proof Construction to Scale zkML, arXiv:2607.08095.
- Governing the Agent-to-Agent Economy of Trust via Progressive Decentralization, arXiv:2501.16606.
- BAML: A decentralized approach to secure, privacy-preserving financial compliance with blockchain hyperledger and federated learning, ResearchGate.
- Blockchain Security Risk Assessment in the Quantum Era, arXiv:2501.11798.